

## ▼ Text Classification

CS 4395.001: Human Language Technologies

Ryan Dimaranan (RTD180003) & Hannah Valena (HCV180000)

### ▼ 1. Dataset

We are using the following text classification dataset from Kaggle: [Twitter Tweets Sentiment Dataset](#). This dataset has 27,500 tweets labeled according to their sentiment. Tweets are labeled as either neutral, positive, or negative.

Our model should be able to predict the sentiment of a given tweet.

```
import pandas as pd

# load dataset
url = 'https://raw.githubusercontent.com/hvalena/nlp-portfolio/main/11-TextClassification/1'
df = pd.read_csv(url, header=0, usecols=['text', 'sentiment'])
print('rows and columns:', df.shape)

# change column types
df['sentiment'] = df.sentiment.astype('category')
df['text'] = df.text.astype('string')

# drop null values
df = df.dropna()

df.head()
```

```
rows and columns: (27481, 2)
```

	text	sentiment
0	I'd have responded, if I were going	neutral
1	Sooo SAD I will miss you here in San Diego!!!	negative
2	my boss is bullying me...	negative
3	what interview! leave me alone	negative
4	Sons of ****, why couldn't they put them on t...	negative

```
import nltk
nltk.download('stopwords', quiet=True)
from nltk.corpus import stopwords
```

! 11s completed at 9:03 PM



```
# perform some preprocessing
stop_words = stopwords.words('english')
df['text'] = df['text'].apply(lambda x: ' '.join([word.lower() for word in x.split() if word not in stop_words])
df.head()
```

	text	sentiment
0	i'd responded, i going	neutral
1	sooo sad i miss san diego!!!	negative
2	boss bullying me...	negative
3	interview! leave alone	negative
4	sons ****, couldn't put releases already bought	negative

```
# split df into train and test
import numpy as np
```


```
np.random.seed(1234)
```


```
i = np.random.rand(len(df)) < 0.8
train = df[i]
test = df[~i]
print("train data size: ", train.shape)
print("test data size: ", test.shape)
```

```
train data size: (21997, 2)
test data size: (5483, 2)
```

```
import seaborn as sb
```

```
# create a graph showing the distribution of the target classes
graph = sb.catplot(x="sentiment", kind="count", data=df)
```





## 2. Sequential Model

```
from keras.preprocessing.text import Tokenizer

vocab_size = 25000

# fit tokenizer on training data
tokenizer = Tokenizer(num_words=vocab_size)
tokenizer.fit_on_texts(train.text)

# convert to numpy matrix
x_train = tokenizer.texts_to_matrix(train.text, mode='tfidf')
x_test = tokenizer.texts_to_matrix(test.text, mode='tfidf')

from sklearn.preprocessing import LabelEncoder

# encode, normalize target values
encoder = LabelEncoder()
encoder.fit(train.sentiment)
y_train = encoder.transform(train.sentiment)
y_test = encoder.transform(test.sentiment)

# check shape
print("train shapes:", x_train.shape, y_train.shape)
print("test shapes:", x_test.shape, y_test.shape)
print("train first 10 labels:", y_train[:10])
print("test first 10 labels:", y_test[:10])

train shapes: (21997, 25000) (21997,)
test shapes: (5483, 25000) (5483,)
train first 10 labels: [1 0 0 0 0 1 2 1 2 0]
test first 10 labels: [1 1 2 1 1 2 0 1 2 1]

from tensorflow.keras import layers, models

# create a sequential model
model = models.Sequential()
model.add(layers.Dense(16, input_dim=vocab_size, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

```
# compile
model.compile(loss='mse',
              optimizer='rmsprop',
              metrics=['accuracy'])
```

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 16)	400016
dense_1 (Dense)	(None, 1)	17
=====		
Total params: 400,033		
Trainable params: 400,033		
Non-trainable params: 0		
=====		

```
# train model
model.fit(x_train, y_train,
          batch_size=128,
          epochs=10,
          verbose=1,
          validation_split=0.1)
```

```
Epoch 1/10
198/198 [=====] - 8s 36ms/step - loss: 0.5437 - accuracy: 0.
Epoch 2/10
198/198 [=====] - 8s 41ms/step - loss: 0.4379 - accuracy: 0.
Epoch 3/10
198/198 [=====] - 6s 31ms/step - loss: 0.4079 - accuracy: 0.
Epoch 4/10
198/198 [=====] - 6s 29ms/step - loss: 0.3908 - accuracy: 0.
Epoch 5/10
198/198 [=====] - 6s 32ms/step - loss: 0.3787 - accuracy: 0.
Epoch 6/10
198/198 [=====] - 6s 31ms/step - loss: 0.3700 - accuracy: 0.
Epoch 7/10
198/198 [=====] - 6s 31ms/step - loss: 0.3632 - accuracy: 0.
Epoch 8/10
198/198 [=====] - 7s 33ms/step - loss: 0.3578 - accuracy: 0.
Epoch 9/10
198/198 [=====] - 6s 33ms/step - loss: 0.3531 - accuracy: 0.
Epoch 10/10
198/198 [=====] - 6s 32ms/step - loss: 0.3492 - accuracy: 0.
```

```
# evaluate
score = model.evaluate(x_test, y_test, batch_size=100, verbose=1)
print('Accuracy: ', score[1])
```

```
print('Accuracy: ', score[1],
```

```
55/55 [=====] - 1s 10ms/step - loss: 0.5057 - accuracy: 0.50  
Accuracy: 0.5013678669929504
```

### 3. RNN & CNN

```
from tensorflow.keras import layers, models
```

```
# RNN
```

```
model = models.Sequential()  
model.add(layers.Embedding(vocab_size, 64))  
model.add(layers.SimpleRNN(64))  
model.add(layers.Dense(1, activation='sigmoid'))  
  
model.compile(optimizer='rmsprop',  
              loss='sparse_categorical_crossentropy',  
              metrics=['accuracy'])
```

```
model.fit(x_train,  
          y_train,  
          epochs=10,  
          batch_size=128,  
          validation_split=0.2)
```

```
score = model.evaluate(x_test, y_test, batch_size=100, verbose=1)  
print('Accuracy: ', score[1])
```

```
from tensorflow.keras import layers, models
```

```
# CNN
```

```
model = models.Sequential()  
model.add(layers.Embedding(25000, 128,))  
model.add(layers.Conv1D(32, 7, activation='relu'))  
model.add(layers.MaxPooling1D(5))  
model.add(layers.Conv1D(32, 7, activation='relu'))  
model.add(layers.GlobalMaxPooling1D())  
model.add(layers.Dense(1))
```

```
model.summary()
```

```
model.compile(optimizer='rmsprop',  
              loss='binary_crossentropy',  
              metrics=['accuracy'])
```

```
model.fit(x_train,
```

```
y_train,  
epochs=10,  
batch_size=128,  
validation_split=0.2)
```

```
score = model.evaluate(x_test, y_test, batch_size=100, verbose=1)  
print('Accuracy: ', score[1])
```

## 4. Embeddings

```
from sklearn.preprocessing import LabelEncoder  
from sklearn.model_selection import train_test_split  
  
# split test and validation  
X_test, X_valid, y_test, y_valid = train_test_split(X, y, test_size=0.2,  
                                                    train_size=0.8,  
                                                    random_state=1234)  
  
# train on all data  
X_train = X  
y_train = y  
  
# labels to integers  
encoder = LabelEncoder()  
encoder.fit(y)  
  
y_valid = encoder.transform(y_valid)  
y_test = encoder.transform(y_test)  
y_train = encoder.transform(y_train)  
  
# vectorize init  
from tensorflow.keras.layers.experimental.preprocessing import TextVectorization  
import tensorflow as tf  
  
vectorizer = TextVectorization(max_tokens=20000, output_sequence_length=200)  
text_ds = tf.data.Dataset.from_tensor_slices(X_train).batch(128)  
vectorizer.adapt(text_ds)  
  
# get vocab  
voc = vectorizer.get_vocabulary()  
word_index = dict(zip(voc, range(len(voc))))  
  
from tensorflow.keras import layers  
from tensorflow import keras  
  
EMBEDDING_DIM = 128  
MAX_SEQUENCE_LENGTH = 200
```

```
# initialize embedding layer
embedding_layer = layers.Embedding(len(word_index) + 1,
                                   EMBEDDING_DIM,
                                   input_length=MAX_SEQUENCE_LENGTH)

import numpy as np

# vectorize data
x_train = vectorizer(np.array([[s] for s in X_train])).numpy()
x_val = vectorizer(np.array([[s] for s in X_valid])).numpy()
x_test = vectorizer(np.array([[s] for s in X_test])).numpy()

y_train = np.array(y_train)
y_val = np.array(y_valid)
y_test = np.array(y_test)

# init model and add layers
int_sequences_input = keras.Input(shape=(None,), dtype="int64")
embedded_sequences = embedding_layer(int_sequences_input)
x = layers.Conv1D(128, 5, activation="relu")(embedded_sequences)
x = layers.MaxPooling1D(5)(x)
x = layers.Conv1D(128, 5, activation="relu")(x)
x = layers.MaxPooling1D(5)(x)
x = layers.Conv1D(128, 5, activation="relu")(x)
x = layers.GlobalMaxPooling1D()(x)
x = layers.Dense(128, activation="relu")(x)
x = layers.Dropout(0.5)(x)

preds = layers.Dense(len(y), activation="softmax")(x)
model = keras.Model(int_sequences_input, preds)
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, None)]	0
embedding (Embedding)	(None, None, 128)	2560128
conv1d (Conv1D)	(None, None, 128)	82048
max_pooling1d (MaxPooling1D)	(None, None, 128)	0
conv1d_1 (Conv1D)	(None, None, 128)	82048
max_pooling1d_1 (MaxPooling1D)	(None, None, 128)	0
conv1d_2 (Conv1D)	(None, None, 128)	82048

global_max_pooling1d (GlobalMaxPooling1D)	(None, 128)	0
dense (Dense)	(None, 128)	16512
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 27480)	3544920

```

=====
Total params: 6,367,704
Trainable params: 6,367,704
Non-trainable params: 0
=====

```

---

```

# compile model and train
model.compile(
    loss="sparse_categorical_crossentropy", optimizer="rmsprop", metrics=["acc"]
)

callback = tf.keras.callbacks.EarlyStopping(monitor='loss', patience=2)

model.fit(x_train, y_train, batch_size=128, epochs=15, validation_data=(x_val, y_val), callbacks=[callback])

```

```

Epoch 1/15
215/215 [=====] - 115s 530ms/step - loss: 1.2898 - acc: 0.37
Epoch 2/15
215/215 [=====] - 112s 521ms/step - loss: 1.0977 - acc: 0.39
Epoch 3/15
215/215 [=====] - 105s 489ms/step - loss: 1.0040 - acc: 0.49
Epoch 4/15
215/215 [=====] - 100s 464ms/step - loss: 0.8088 - acc: 0.66
Epoch 5/15
215/215 [=====] - 104s 485ms/step - loss: 0.6864 - acc: 0.74
Epoch 6/15
215/215 [=====] - 99s 461ms/step - loss: 0.5591 - acc: 0.788
Epoch 7/15
215/215 [=====] - 101s 469ms/step - loss: 0.4311 - acc: 0.84
Epoch 8/15
215/215 [=====] - 99s 461ms/step - loss: 0.3336 - acc: 0.885
Epoch 9/15
215/215 [=====] - 100s 463ms/step - loss: 0.2482 - acc: 0.91
Epoch 10/15
215/215 [=====] - 99s 459ms/step - loss: 0.1807 - acc: 0.940
Epoch 11/15
215/215 [=====] - 98s 455ms/step - loss: 0.1304 - acc: 0.959
Epoch 12/15
215/215 [=====] - 95s 444ms/step - loss: 0.0966 - acc: 0.969
Epoch 13/15
215/215 [=====] - 97s 450ms/step - loss: 0.0721 - acc: 0.978
Epoch 14/15
215/215 [=====] - 96s 445ms/step - loss: 0.0583 - acc: 0.982
Epoch 15/15

```



```
-- -- --
215/215 [=====] - 96s 448ms/step - loss: 0.0451 - acc: 0.986
<keras.callbacks.History at 0x7f4769b86190>
```

```
# test and evaluate model
```

```
score = model.evaluate(x_test, y_test, batch_size=128, verbose=1)
```

```
print('Accuracy: ', score[1])
```

```
172/172 [=====] - 30s 172ms/step - loss: 0.3378 - acc: 0.879
Accuracy: 0.8789574503898621
```

```
!wget --no-check-certificate http://nlp.stanford.edu/data/glove.6B.zip
```

```
!unzip glove.6B.zip
```

```
import os
```

```
import numpy as np
```

```
# get glove embeddings
```

```
path_to_glove_file = os.path.join(
```

```
    "glove.6B.200d.txt"
```

```
)
```

```
num_tokens = len(voc) + 2
```

```
embedding_dim = 200
```

```
hits = 0
```

```
misses = 0
```

```
embeddings_index = {}
```

```
with open(path_to_glove_file) as f:
```

```
    for line in f:
```

```
        word, coefs = line.split(maxsplit=1)
```

```
        coefs = np.fromstring(coefs, "f", sep=" ")
```

```
        embeddings_index[word] = coefs
```

```
embedding_matrix = np.zeros((num_tokens, embedding_dim))
```

```
for word, i in word_index.items():
```

```
    embedding_vector = embeddings_index.get(word)
```

```
    if embedding_vector is not None:
```

```
        embedding_matrix[i] = embedding_vector
```

```
        hits += 1
```

```
    else:
```

```
        misses += 1
```

```
from tensorflow.keras.layers import Embedding
```

```
# initialize embedding layer with pretrained embeddings
```

```
embedding_layer = Embedding(
```

```
num_tokens,
embedding_dim,
embeddings_initializer=keras.initializers.Constant(embedding_matrix),
trainable=False,
)

# compile and train with early stopping
model.compile(
    loss="sparse_categorical_crossentropy", optimizer="rmsprop", metrics=["acc"]
)

callback = tf.keras.callbacks.EarlyStopping(monitor='loss', patience=2)

model.fit(x_train, y_train, batch_size=128, epochs=15, validation_data=(x_val, y_val), cal:

Epoch 1/15
215/215 [=====] - 87s 399ms/step - loss: 0.4466 - acc: 0.826
Epoch 2/15
215/215 [=====] - 88s 409ms/step - loss: 0.3826 - acc: 0.851
Epoch 3/15
215/215 [=====] - 85s 398ms/step - loss: 0.3388 - acc: 0.873
Epoch 4/15
215/215 [=====] - 88s 408ms/step - loss: 0.3086 - acc: 0.885
Epoch 5/15
215/215 [=====] - 85s 398ms/step - loss: 0.2750 - acc: 0.899
Epoch 6/15
215/215 [=====] - 88s 408ms/step - loss: 0.2555 - acc: 0.906
Epoch 7/15
215/215 [=====] - 93s 435ms/step - loss: 0.2287 - acc: 0.918
Epoch 8/15
215/215 [=====] - 86s 401ms/step - loss: 0.2024 - acc: 0.928
Epoch 9/15
215/215 [=====] - 88s 407ms/step - loss: 0.1913 - acc: 0.932
Epoch 10/15
215/215 [=====] - 93s 432ms/step - loss: 0.1808 - acc: 0.937
Epoch 11/15
215/215 [=====] - 87s 403ms/step - loss: 0.1685 - acc: 0.941
Epoch 12/15
215/215 [=====] - 85s 398ms/step - loss: 0.1513 - acc: 0.947
Epoch 13/15
215/215 [=====] - 87s 402ms/step - loss: 0.1521 - acc: 0.948
Epoch 14/15
215/215 [=====] - 85s 397ms/step - loss: 0.1369 - acc: 0.953
Epoch 15/15
215/215 [=====] - 87s 405ms/step - loss: 0.1313 - acc: 0.953
<keras.callbacks.History at 0x7fdcf0cb9940>

# evaluate model
score = model.evaluate(x_test, y_test, batch_size=128, verbose=1)
print('Accuracy: ', score[1])

172/172 [=====] - 29s 167ms/step - loss: 0.2370 - acc: 0.913
Accuracy: 0.9137099981307983
```